


Bobcat at a glance


Accuracy

Bobcat Same base model, zero-shot



Agreement vs time per case

TypeSafe's 20 published workflow examples



One decision: 50 ms p50

512 tokens, 8 candidates, one H100, FP8, server-side

Accuracy: Bobcat vs Qwen3.8-27B zero-shot on identical inputs. Right: agreement with the mean of GPT-6 Astra and Claude Fable 5.1 on 329 questions; Jev, Opus 5 and Sol are TypeSafe's published answers and times (client-side, network included); Bobcat was run on one L40S, server-side, step requests in sequence.