



■ HTTP p50    ● HTTP p95    - - - W1 targets (p50 150 ms, p95 300 ms); they apply to W1 only

Latency per request profile



Throughput, W4 (labels: p50 ms)



1 B300 GPU, bfloat16, merged LoRA; warm; in-process HTTP client (no network).